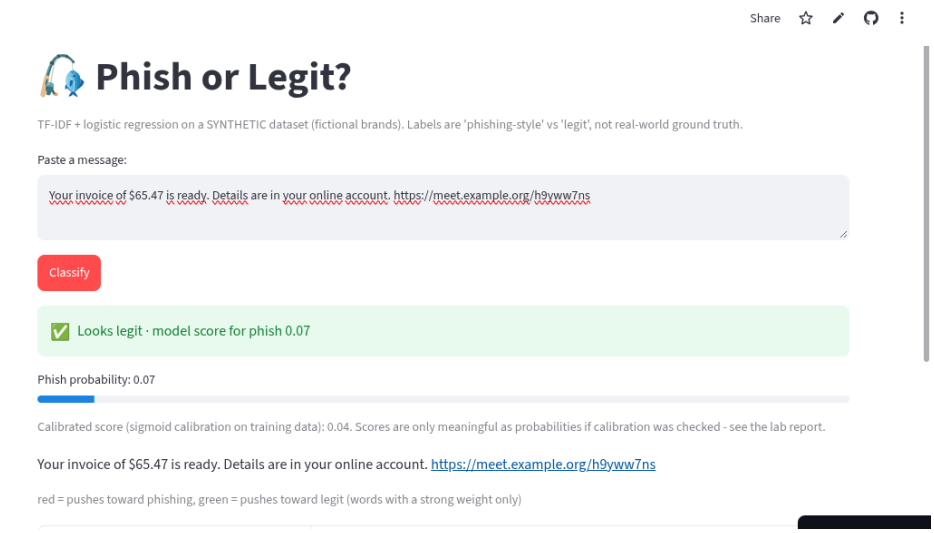
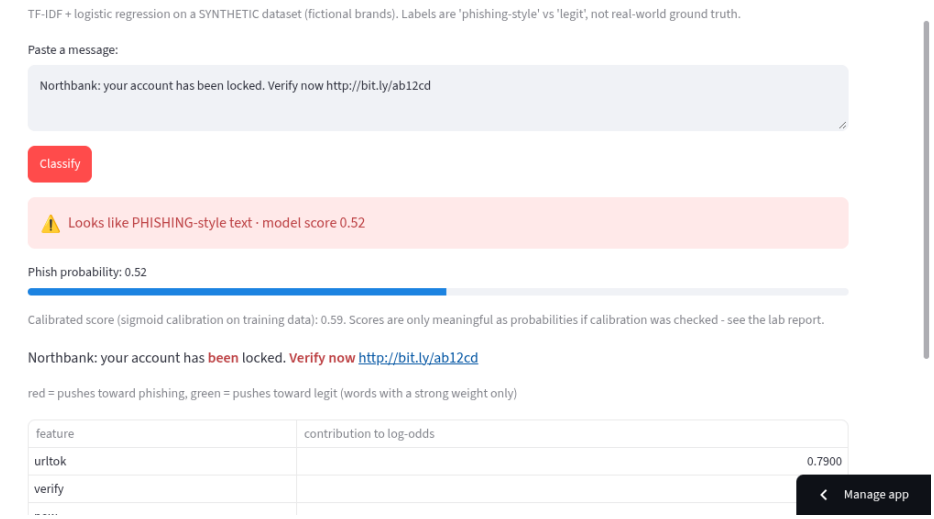
Lab 04 : Phish or Legit? Report



Legit example



Phishing example

Purpose

This project builds a text classifier that decides whether a short message is **phishing-style** or **legit**, using a synthetic dataset of fictional-brand messages. It is an educational lab (not a production

detector): the goal is to show the full pipeline - data, features, model, evaluation, calibration, and explainability - and why accuracy alone is misleading on imbalanced data.

Dataset

- `data/messages.csv` - 2000 synthetic messages, 20% phishing / 80% legit.
- Fictional brands (Northbank, ParcelGo, Civic Savings, …); URLs, money amounts and numbers are normalised to placeholder tokens (`urltok`, `moneytok`, `numtok`).
- Stratified split: 1500 train / 500 test, keeping the 20% phishing rate in both.

Methodology

- Features: TF-IDF (word unigrams+bigrams, sublinear tf, min_df=2).
- Models compared: Naive Bayes, Logistic Regression, NB + sigmoid calibration, LR + char n-grams (extension).
- All fitting happens inside a single scikit-learn `Pipeline` on the **train** split only, to avoid data leakage.
- Split is seeded (SEED=42) and stratified.

Results (test set, 500 messages)

Model	Precision	Recall	F1	Accuracy	Brier	ECE
Naive Bayes	0.611	0.800	0.693	0.858	0.104	0.089
Logistic Regression	0.747	0.710	0.728	0.894	0.071	0.035
NB + calibration	0.694	0.750	0.721	0.884	0.086	0.030
LR + char n-grams	0.777	0.730	0.753	0.904	0.064	0.050

Majority-class baseline accuracy (always “legit”): **80.0%** - which is why accuracy alone is misleading; F1 / precision / recall on the phish class matter.

Confusion matrix (Logistic Regression)

Rows = truth, columns = prediction (legit, phish):

```
[[376  24]
 [ 29  71]]
```

Decision threshold trade-off (Logistic Regression)

Threshold	Precision	Recall	F1
0.20	0.604	0.930	0.732
0.35	0.672	0.800	0.731
0.50	0.747	0.710	0.728
0.70	0.893	0.500	0.641
0.90	1.000	0.090	0.165

Lowering the threshold catches more phishing (higher recall) at the cost of more false alarms; raising it does the opposite.

Strongest features (Logistic Regression)

- Phish: `urltok` (4.83), `immediately` (3.98), `notice` (3.60), `immediately urltok` (3.17), `alert` (2.82), `can urltok` (2.81)
- Legit: `reminder` (-2.77), `hey` (-2.52), `numtok` (-2.28), `fyi` (-2.01), `order` (-1.63), `order numtok` (-1.63)

Failure cases

- False negatives (phish → legit): 29
- False positives (legit → phish): 24

See `results/failures.txt` for examples. Common failure: phishing messages that mimic legit tone (“Your Civic Savings account has a new notice…”) slip through.

Calibration

Brier score and ECE show how well scores behave as probabilities. Logistic Regression (ECE 0.035) is better calibrated than Naive Bayes (ECE 0.089); sigmoid calibration improves NB further. The app shows

both the raw model score and a sigmoid-calibrated score. Scores are only meaningful as probabilities if calibration was checked - see the reliability diagram in `results/phish_results.png`.

Explainability in the app

For each classified message, the app: 1. Shows the predicted label and score (plus calibrated score). 2. Highlights words that strongly push toward phish (red) or legit (green). 3. Lists the top contributing features to the log-odds.

How to run

```
pip install -r requirements.txt
streamlit run app.py          # UI
python -m src.evaluate        # metrics + plots in results/
```

Files

- `app.py` - Streamlit UI
- `src/pipeline.py` - data loading, model building, prediction, explanations
- `src/evaluate.py` - metrics, confusion matrix, calibration, failure cases
- `data/messages.csv` - synthetic dataset
- `data/make_data.py` - dataset generator
- `results/` - metrics.json, failures.txt, phish_results.png

Limitations

- Synthetic data with fictional brands; not real-world ground truth.
- Small dataset; models may not generalise to real phishing.
- Word-level highlighting only covers tokens with a strong weight (|w| > 0.3).
- Calibrated scores assume the training distribution matches deployment. “””